

Paid to Hold the Bag

Liquidity provision is a high-VIX premium — at the daily horizon and at the opening auction — and that is exactly why it is hard to size

Gator Quant Hacks 2026 · Track 03 Systematic Trading · October 2026 · All numbers: `python run_all.py --oos` in the public repo · Hypotheses pre-registered in HYPOTHESIS.md (commit 0d0093d, 2026-10-03 10:54 UTC) HYPOTHESIS_B.md (commit fbd355a, 11:53 UTC) and HYPOTHESIS_C.md (commit 6067ea1, 18:14 UTC), each before its first backtest

1. Summary

Short-term reversal is the return to supplying liquidity. Before looking at any result we hypothesised that this premium is larger when liquidity-provider capital is scarce, proxied by yesterday's VIX, and tested the one mechanism on two instruments, then tried to break our own conclusion. **Module A** (36 years of Kenneth French daily reversal portfolios): the expected reversal return rises with lagged VIX (Newey–West $t = 2.4$); since 2004 the large-cap premium exists *only* in the top VIX quintile (14 bps/day when $VIX > 24$, ≈ 0 otherwise; 52 bps/day on the 44 such held-out days, nothing on the other 455). The pre-registered rule that *levers* into high VIX fails: the book's volatility rises $3.6\times$ across the same quintiles, so more dollars buy the same Sharpe. **Module B** takes the lesson to the opening auction: fading overnight gaps in S&P 500 names open-to-close earns 18 bps/day gross ($t = 12$, gross Sharpe 3.3), 35 bps/day in the top VIX quintile versus 12 elsewhere. At 5 bps per dollar traded it is a market maker's premium — always-on nets nothing — and the pre-registered rule, *participate* only when $VIX_{t-1} \geq 20$, turns -0.26 Sharpe into $+0.12$ in-sample and -1.83 into $+0.24$ on the held-out two years, when the top-VIX state paid 56 bps/day and every other state zero. The premium has decayed (33 bps/day before 2013, 8 after). **Module C** tried to beat VIX with a sharper state variable, the variance risk premium (price of risk) separated from realised variance (quantity of risk), and was **rejected**: the premium loads on realised variance and VRP rules lose to constant exposure in and out of sample. What survives for a non-market-maker is a constant-notional large-cap reversal book (net Sharpe ≈ 0.4 , break-even 14 bps, \$100M–\$500M capacity) and, for a desk with 2.5 bps auction execution, a VIX-gated gap fade (Sharpe ≈ 1). All net of costs; all 80 variants reported.

2. Economic hypothesis (written before results)

Who is on the other side. Someone who must trade *now*: a fund meeting redemptions, a forced deleverager, an investor reacting to overnight news with a market-on-open order. They pay, through a temporarily displaced price, for a counterparty; the price then reverts. A reversal book, daily or intraday, is a systematic way of being that counterparty. **Why it should depend on VIX.** Liquidity providers hold inventory risk; when volatility is high their VaR limits bind, margins rise and capital is withdrawn, so the price of immediacy rises (Grossman–Miller 1988; Brunnermeier–Pedersen 2009; Nagel 2012). The edge persists because the capital that would compete it away is what is scarce when the premium is highest; at the open the same logic applies to the overnight gap (Berkman et al. 2012; Lou, Polk & Skouras 2019).

Pre-registered predictions. Module A — P1: slope of the daily reversal return on VIX_{t-1} positive, NW $t > 2$; P2: mean rises across lagged-VIX quintiles; P3: $w_t = \min(VIX_{t-1}/20, 3)$ beats $w = 1$ net; P4: alpha survives FF5 + momentum; P5: neighbouring parameters agree. Module B, written after A's result and before any stock data was downloaded — B1: gross gap-fade spread positive, $t > 2$; B2: slope on VIX_{t-1} positive; B3: at 5 bps, a $VIX \geq 20$ participation rule beats always-on, with positive net mean; B4: both legs earn it (long leg only would suggest survivorship); B5: plateau over thresholds $\{15, 20, 25\}$ and $\{\text{quintile, decile}\}$. Module C, written after both — C1: in a joint regression the premium loads on VRP_{t-1} , not RV_{t-1} ; C2: within realised-variance terciles, Sharpe rises with VRP; C3: a VRP-participation rule

beats constant exposure net; C4: VIX-term-structure backwardation beats constant, 2008+. P3, half of B3, and C1–C3 failed; kill criteria were written for each.

3. Data and universe

All public, no keys. **A:** Kenneth French Data Library (CRSP-based, point-in-time, through Aug 2026): the daily **ST_Rev** factor (long the bottom 30% on day -20 to -1 return, short the top 30%, **re-formed daily**); the daily 6 size $\times$ prior-return portfolios, from which we build the **big-cap reversal leg** BIG Lo – BIG Hi; FF5 + momentum. **B:** daily adjusted open and close for the current S&P 500 constituents (Wikipedia list snapshotted and committed; prices from Yahoo via yfinance, cached, not committed). Each stock enters **only from its index-addition date** (326 eligible names/day in-sample, 480 recently); removed names are absent, so the universe has survivorship bias — disclosed, and tested by B4. **All:** FRED CBOE VIX close (and 3-month VIX, from Dec 2007, for Module C), aligned by a strict as-of merge so the position on day t sees only closes dated before t (unit-tested).

Samples: A 1990-01-02 $\rightarrow$ 2026-08-31 (8,732 in-sample days); B 2005-01-03 $\rightarrow$ 2026-08-31 (4,949). **Held out for both: 2024-09-01 $\rightarrow$ 2026-08-31** (the 2-year cap binds; 500 days), evaluated once per module, each in its own commit after the in-sample commit.

4. Methodology

Module A book and rule. One unit is \$1 long losers / \$1 short winners. $w_t = \min(\text{VIX}_{t-1}/20, 3)$: 20 is the textbook long-run VIX, so $w = 1$ at VIX 20 and the rule is comparable to constant exposure; 3 is a leverage limit; nothing is fitted. **Costs:** a 20-day formation window implies 12–15% of each side turns over daily; we assume $\tau = 14\%$ one-way, dollars traded $= 2w_t\tau + 2|\Delta w_t|$, and cost per dollar $c_t = c_{\text{base}} \times \text{VIX}_{t-1}/20$ (spreads widen with VIX), $c_{\text{base}} = 5$ bps large-cap, 10 bps all-cap; flat, $\times 2$ and break-even costs also reported.

Module B book and rule. For each stock, $\text{gap} = \text{Open}_t/\text{Close}_{t-1} - 1$ minus the cross-sectional median gap; sort into quintiles at the open; \$1 long the most-gapped-down quintile, \$1 short the most-gapped-up, equal-weighted, **entered at the official open and closed at the official close** — never an overnight position. Rule: on iff $\text{VIX}_{t-1} \geq 20$, benchmark always-on. Four dollars are traded per day per \$1-per-side, so the toll is $4 \times 5 \text{ bps} \times \text{VIX}_{t-1}/20 \approx 20 \text{ bps/day}$; $c_{\text{base}} \in \{2.5, 10\}$ and the break-even are reported. Stock-days with $|\text{gap}| > 20\%$ (spin-off artefacts) are excluded using open information only; without the filter results are identical to 0.1 bps.

Tests and budget. Slopes by OLS with Newey–West (10 lags); rule comparisons by a spanning regression of rule on benchmark (HAC t on the intercept); A's alpha by FF5 + Mom. Grids: A, normaliser $\{15, 20, 25\} \times \text{cap} \{2, 3, 4\} \times \text{universe} \times \{\text{linear, regime}\} = 36$, plus 2 post-hoc rules (\$5.4) = 38 trials entering the deflated Sharpe (Bailey & López de Prado 2014); B, threshold $\{\text{off}, 15, 20, 25\} \times \{\text{quintile, decile}\} \times c_{\text{base}} \{2.5, 5, 10\} = 24$; C, realised-variance window $\{21, 63\} \times \text{rule} \{\text{VRP} > \text{trailing median, VRP} > 0, \text{VIX} > \text{VIX3M}\} \times \text{book} \{\text{big-cap, ST_Rev, gap fade}\} = 18$, on the same books and costs. **80 variants, all reported.**

5. Results

5.1 Module A: the premium depends on yesterday's VIX (P1 passes, P2 passes in shape; both replicate)

Mean next-day return, bps	slope / VIX pt	NW t	Q1	Q2	Q3	Q4	Q5 (VIX > 24)
Big-cap leg, in-sample 1990–2024	0.89	2.42	0.4	1.8	−1.0	2.2	16.1 $\pm$ 9.1
ST_Rev, in-sample	0.78	2.45	5.2	5.4	2.8	5.7	17.6 $\pm$ 7.6
Big-cap leg, out-of-sample	2.94	2.50	–	6.9	−4.3	−13.1	52.4 $\pm$ 50.7
ST_Rev, out-of-sample	2.83	3.10	–	3.1	−1.9	−8.2	47.3 $\pm$ 44.3
Big-cap leg, 2004–2024 only	0.82	1.66	−3.3	1.2	−4.5	0.2	14.0

$\pm$ is a 95% CI. 1,746 days per in-sample quintile; edges 13.2 / 16.1 / 19.6 / 24.3; out-of-sample binned on those edges (44 days in Q5; Q1 has one day, omitted). Constant-book net Sharpe 1.02 (big-cap) / 2.25 (ST_Rev) in 1990–2003, –0.02 for both in 2004–2024.

P1 passes and replicates with the same sign and a larger slope (figure in the appendix). P2 is not monotone: flat across quintiles 1–4, then a jump. For large caps the premium is *only* in quintile 5, and the sub-period row shows why: after decimalisation and electronic market making the unconditional premium went to zero, as the hypothesis says it should when liquidity-provider capital becomes abundant; what survives is the part earned when that capital withdraws.

5.2 Module A's rule: more dollars, not more Sharpe (P3 fails, P4 passes)

Net of costs	Constant, IS	VIX-timed, IS	Timed, flat costs	Timed, costs $\times 2$	Constant, OOS	Timed, OOS
Big-cap: ann. return / vol	6.4% / 17.2%	10.9% / 30.3%	11.5% / 30.3%	5.6% / 30.3%	2.7% / 20.5%	8.0% / 23.8%
Big-cap: net Sharpe (gross)	0.37 (0.57)	0.36 (0.53)	0.38	0.18	0.13 (0.29)	0.34 (0.54)
ST_Rev: net Sharpe (gross)	0.82 (1.30)	0.51 (0.93)	0.56	0.10	0.04 (0.45)	0.21 (0.70)
Break-even c_{base} : big-cap / ST_Rev	14.3 / 26.9 bps	15.3 / 22.4 bps				

The timed rule lifts big-cap gross return from 9.9% to 16.2% at the same mean exposure, as predicted, but volatility goes from 17% to 30%: the book's own volatility scales with VIX, so sizing linearly in VIX scales risk quadratically. Spanning alpha of timed on constant: +0.24 bps/day ($t = 0.3$) big-cap, **–2.35 bps/day ($t = -3.0$)** ST_Rev, whose calm-state premium the rule discards. P3 fails. P4 passes: FF5 + Mom alpha 5.8%/yr ($t = 2.1$) constant big-cap, 11.1%/yr ($t = 4.8$) ST_Rev; market beta 0.2. The grid is a plateau (0.34–0.39; P5 passes); luck alone would give a best-of-38 Sharpe of 0.24. Out of sample the timed rule beat constant, but with spanning $t = 1.1$ after failing on 35 years we do not count it: the premium arrived on 44 high-VIX days and a rule long VIX happened to be large on them (equity curves in the appendix).


5.3 Module B: the opening auction pays the same premium, in the same states (B1, B2, B4, B5 pass; B3 half)

Q1 – Q5 gap-fade spread, open $\rightarrow$ close	gross bps/day (NW t)	slope per VIX pt (t)	VIX Q1 ... Q5, bps	net Sharpe @5 bps: always-on / gated	spanning t	break-even c_{base}	@2.5 bps: always / gated	days on
In-sample 2005–2024	17.8 (12.1)	1.39 (3.8)	12 $\cdot$ 12 $\cdot$ 11 $\cdot$ 19 $\cdot$ 35 $\pm$ 10	–0.26 / +0.12	2.6	4.6 / 5.3 bps	1.51 / 1.12	33%
2005–2012	32.6 (12.3)	1.69 (3.7)	24 $\rightarrow$ 59	1.84 / 1.19		7.6 / 7.6		
2013–2024	7.6 (5.6)	0.26 (0.8)	5 $\rightarrow$ 11	–2.06 / –1.00		2.2 / 2.3		
Out-of-sample 2024-09 $\rightarrow$ 2026-08	6.7 (1.4)	3.98 (3.3)	– $\cdot$ 0 $\cdot$ 0 $\cdot$ 4 $\cdot$ 56 $\pm$ 43	–1.83 / +0.24	2.0	1.8 / 5.9	–0.40 / 0.92	24%

Equal-weighted quintiles on the gap relative to the median gap; \$1 per side; cost $4 \times c_{\text{base}} \times \text{VIX}_{t-1}/20$ per day. OOS VIX bins use in-sample edges (Q1 empty). Legs, excess over the universe's open $\rightarrow$ close return: long gap-down +9.2 bps ($t = 11.8$), short gap-up –8.6 bps ($t = -11.2$) in-sample; +4.4 / –2.3 out of sample. Decile sort in appendix A2.

Gross, this is a 3-Sharpe machine, and it is the mechanism of §5.1 at a shorter horizon: 35 bps/day in the top VIX quintile against 12 in the rest, both legs symmetric (B4 passes; survivorship is not the story), slope $t = 3.8$. Net, at 5 bps per dollar traded, the ≈ 20 bps/day toll eats the unconditional premium (always-on –0.26), and the participation rule pays the toll only when the premium exceeds it, lifting Sharpe to +0.12 (spanning $t = 2.6$) — but its net mean is not significantly positive ($t = 0.5$): **at 5 bps this is a market maker's premium**. Two facts matter more. The premium is decaying: 33 bps/day and a 7.6 bps break-even before 2013, 8 bps and 2.2 bps since, as electronic market making in the auctions caught up; a 2.5 bps executor earned Sharpe 1.5 over the full sample and roughly zero since 2013. And the held-out two years are the sharpest state-dependence we have seen anywhere: unconditional gross fell to 6.7 bps,


the VIX slope rose to 4 bps per point ($t = 3.3$), the 50 days with $VIX_{t-1} > 24$ paid **56 bps/day** and every other state paid zero. Always-on lost 30% net; the gated book made 2.5% with a 9% drawdown, Sharpe 0.9 at 2.5 bps (equity curves in the appendix).



5.4 Post-hoc, disclosed: sizing Module A by risk instead (2 extra trials)

Sizing by expected return per unit of variance — $w_t \propto \max(E[r_t | VIX_{t-1}], 0)/\hat{\sigma}_t^2$, with $E[r | VIX]$ from an expanding regression on days before t and $\hat{\sigma}$ the trailing 63-day vol — against a vol-targeted book (Moreira-Muir 2017) and constant exposure on the same days (big-cap, 1995–2024, net Sharpe / max DD / worst month): constant **0.31** / -48% / -17%; vol-targeted **0.01** / -70% / -14%; mean-variance **0.19** / -74% / -53%. Vol targeting, the default risk tool, **destroys the strategy**: it de-risks exactly when the premium is paid. The mean-variance rule's -53% month is March 2020: the forecast said "premium high", the rule went to the $3\times$ cap, and reversal kept losing. This is why Module B pre-registered participation rather than leverage, and why Module C asked whether any state variable can select the premium without the volatility.

5.5 Module C: is it the price of risk or the quantity of risk? (C1-C3 fail; C4 does not replicate)



Net Sharpe by switching rule	Constant / always-on	VIX rule (A: regime; B: $VIX \geq 20$)	VRP > trailing median	VRP > 0	VIX > VIX3M (2008+)	Constant, 2008+
Big-cap leg, IS → OOS	0.37 → 0.13	0.39 → -0.05	0.03 → 0.02	0.13 → -0.07	0.26 → 0.36	0.05
ST_Rev, IS → OOS	0.82 → 0.04	0.47 → -0.17	-0.13 → -0.33	0.47 → -0.26	0.11 → 0.13	0.03
Gap fade, IS → OOS	-0.26 → -1.83	0.12 → 0.24	-0.28 → -0.87	-0.36 → -1.57	0.65 → -0.18	-0.78

C1 in-sample, r_t on VRP_{t-1} and RV_{t-1} (21-day): $t_{VRP} / t_{RV} = 0.7 / 1.9$ big-cap, $0.4 / 1.9$ ST_Rev, $2.8 / 3.1$ gap fade. Spanning alpha of the VRP-median rule on constant: $t = -2.0$ (big-cap), -5.7 (ST_Rev). Backwardation on the gap fade: spanning $t = 4.3$ in-sample on 10% of days; -0.18 out of sample against 0.24 for $VIX \geq 20$.

VIX bundles the *quantity* of risk a liquidity provider must hold (realised variance) with the *price* paid for holding it (the variance risk premium, $VIX^2 - RV$). If the premium were a price-of-risk effect, VRP would select the good days without the volatile ones. It does not: the premium loads on realised variance, not VRP; the big-cap Sharpe is a function of the realised-variance row of the double sort (≈ 1.0 in the top

tercile, ≈ 0 elsewhere), not the VRP column; and every VRP rule loses to constant exposure in-sample and out. The one proxy that helped in-sample — term-structure backwardation, a 10%-of-days stress flag — was 1 of 18 variants, helped big caps only modestly (spanning $t = 1.3$) and failed on the gap fade out of sample. Reading: liquidity-provider capital is constrained by **realised** turbulence (VaR limits and margins are set on realised vol), so the premium arrives with the volatility and no public state variable we found separates them — the strongest support yet for the constant-notional conclusion.

6. Risk management

What the books are. A: long/short equity, residual market beta 0.2, skew +1.2, excess kurtosis 23 at constant notional (max drawdown -48%, worst month -17%); the timed version has skew +4 and kurtosis 127, a P&L made of a few enormous days (-77% / -34%). B: dollar-neutral, equal-weighted, flat at every close, so no overnight, gap or financing risk; its risk is one session's cross-sectional dispersion (worst day -6.6%, 9 Nov 2020, the vaccine rotation: those gaps were information). **Limits:** gross $\leq 2\times$ capital; per-name $\leq 2\%$ of gross with a mega-cap concentration cap; sector-neutral legs; a futures hedge taking beta to zero. **De-risking rule, set in advance: none that depends on realised vol**, since §5.4–5.5 show the premium is the realised-vol state. *Participation* is the risk tool instead: B is off two-thirds of the time and, when on, sized so the pre-registered worst month at $1\times$ (A -17%; B -16%) and drawdown (A -48%; B gated -79%, a slow bleed 2010–2023 while the premium thinned) are survivable — a 5%-per-month risk budget implies $0.3\times$ capital for either. **Tail and regime:** the book loses when the hurried trader knew something (2009, 2020); those are the risks being insured. In quintiles 1–4 the modern premium is zero, so a non-market-maker is paid nothing for the risk it runs there.

7. Liquidity and capacity

A turns over $70\times$ its capital a year. One-way cost = half-spread + square-root impact, $c(K) = 1.5 \text{ bps} + \sigma_{\text{name}} \sqrt{\text{participation}}$, ≈ 500 names, value-weighted ADV $\approx \$200\text{M}$, single-name vol 2%/day at VIX 20 (src/gqh/capacity.py): $\$10\text{M} \rightarrow 2.6 \text{ bps}$, Sharpe 0.47; $\$100\text{M} \rightarrow 4.8 \text{ bps}$, 0.38; $\$500\text{M} \rightarrow 9.0 \text{ bps}$, 0.21; $\$1.5\text{B} \rightarrow 14.3 \text{ bps}$, 0.00. Our 5 bps is about $\$100\text{M}$; the edge is half gone near $\$500\text{M}$. B trades only in the two auctions, at one price with no spread to cross; its constraint is the open, which clears only ~ 1 –2% of a name's daily volume. With ~ 80 names per leg, $\$25\text{M}$ of book is $\sim \$150\text{k}$ per name, about 5% of a $\$3\text{M}$ opening auction in a $\$200\text{M}$ -ADV name; $\$100\text{M}$ would be $\sim 20\%$, where impact sets the cost. B is a desk-scale trade — tens of millions, not hundreds; 5 bps is our headline, 2.5 bps an execution desk's number. A market maker who *earns* the half-spread faces a negative first term in both books and the gross Sharpes (0.57 and 3.3) are on offer, which is why both premia belong to market makers.

8. Limitations and next steps

(i) The French portfolios are a test instrument, not a product; B's universe is survivorship-biased (removed names missing), though the symmetric legs say the result is not drift. (ii) A's turnover and both cost models are assumed and stressed, not observed. (iii) B assumes execution at the official open from orders placed on pre-market prints; that slippage sits in the cost assumption, and part of the premium may live inside the first minutes. (iv) The top-quintile effect rests on $\sim 1,750$ in-sample and 44–50 held-out days: the sign replicated with 2 – $4\times$ the slope, but the 52–56 bps carry $\pm 45 \text{ bps}$ intervals. (v) "VIX > 24" was deliberately *not* turned into a rule; Yahoo data can be revised, so the download is cached and the constituent file frozen. (vi) Module C used the two daily public price-of-risk proxies; dealer-balance-sheet measures are quarterly and untried. **Next:** signed auction imbalance from exchange feeds; a Q5-only book on a sealed window before believing 50 bps; intermediary leverage at the monthly horizon.

References (outside the page limit)

Almgren, R., Thum, C., Hauptmann, E. & Li, H. (2005). Direct Estimation of Equity Market Impact. *Risk* 18. · Bailey, D. H. & López de Prado, M. (2014). The Deflated Sharpe Ratio. *Journal of Portfolio Management* 40(5). · Berkman, H., Koch, P. D., Tuttle, L. & Zhang, Y. J. (2012). Paying Attention: Overnight Returns and the Hidden Cost of Buying at the Open. *Journal of Financial and Quantitative Analysis* 47(4). · Brunnermeier, M. K. & Pedersen, L. H. (2009). Market Liquidity and Funding Liquidity. *Review of Financial Studies* 22(6). · Chordia, T., Subrahmanyam, A. & Tong, Q. (2014). Have Capital Market Anomalies Attenuated in the Recent Era of High Liquidity and Trading Activity? *Journal of Accounting and Economics* 58. · Fama, E. F. & French, K. R. (2015). A Five-Factor Asset Pricing Model. *Journal of Financial Economics* 116. · Grossman, S. J. & Miller, M. H. (1988). Liquidity and Market Structure. *Journal of Finance* 43(3). · Heston, S. L., Korajczyk, R. A. & Sadka, R. (2010). Intraday Patterns in the Cross-section of Stock Returns. *Journal of Finance* 65(4). · Jegadeesh, N. (1990). Evidence of Predictable Behavior of Security Returns. *Journal of Finance* 45(3). · Lou, D., Polk, C. & Skouras, S. (2019). A Tug of War: Overnight Versus Intraday Expected Returns. *Journal of Financial Economics* 134(1). · Moreira, A. & Muir, T. (2017). Volatility-Managed Portfolios. *Journal of Finance* 72(4). · Nagel, S. (2012). Evaporating Liquidity. *Review of Financial Studies* 25(7). · Data: Kenneth R. French Data Library (CRSP-based); FRED, Federal Reserve Bank of St. Louis, series VIXCLS; Yahoo Finance daily prices via yfinance; Wikipedia, List of S&P 500 companies (snapshot 2026-10-03).

Appendix (optional, outside the page limit)

A1. Module A, all 36 pre-registered variants, net Sharpe in-sample. Big-cap, linear rule, rows cap {2,3,4} × cols normaliser {15,20,25}: 0.39 0.37 0.36 / 0.37 0.36 0.35 / 0.36 0.35 0.34; regime rule 0.39. ST_Rev: 0.62 0.57 0.54 / 0.55 0.51 0.49 / 0.51 0.48 0.48; regime rule 0.47. results/tables/sensitivity_grid.csv.

A2. Module B, all 24 pre-registered variants, net Sharpe in-sample (rows c_{base} 2.5 / 5 / 10 bps; cols always-on, VIX ≥ 15, 20, 25). Quintile: 1.51 1.19 1.12 0.87 / -0.26 -0.20 0.12 0.17 / -3.80 -2.97 -1.88 -1.22. Decile: 2.32 1.88 1.65 1.27 / 0.97 0.81 0.89 0.75 / -1.77 -1.35 -0.67 -0.33. Out of sample, quintile: -0.40 -0.29 0.92 1.04 / -1.83 -1.64 0.24 0.68 / -4.73 -4.33 -1.13 -0.07. results/tables/gapfade_sensitivity_grid*.csv, results/REPORT_B.md.

A3. Module C, all 18 pre-registered variants, net Sharpe in-sample → out of sample (RV window 21 / 63; rules VRP > median, VRP > 0, backwardation). Big-cap: 0.03 → 0.02 / 0.16 → -0.23, 0.13 → -0.07 / 0.20 → 0.30, 0.26 → 0.36 (both windows). ST_Rev: -0.13 → -0.33 / 0.04 → -0.66, 0.47 → -0.26 / 0.55 → 0.04, 0.11 → 0.13. Gap fade: -0.28 → -0.87 / 0.08 → -1.07, -0.36 → -1.57 / 0.07 → -1.68, 0.65 → -0.18. results/tables/pricing_variant_grid*.csv, results/REPORT_C.md.

A4. Big-cap leg, net return by year, timed / constant (%). 1990–99: 6/3, 13/14, 9/13, 3/4, 15/24, 15/25, 18/21, 11/12, 34/17, -5/-3. 2000–09: 34/23, 69/50, 56/28, 16/9, 0/0, -8/-13, -3/-5, -5/-9, 61/31, -19/0. 2010–24: 4/2, 1/-2, -11/-11, 1/2, 2/0, -2/-7, 7/12, -3/-5, -2/-4, 11/13, -22/-5, -7/-12, 0/-3, -15/-12, 2/-4. **Gap fade, net by year, gated / always-on (%):** 2005–09: 0/17, -2/35, 41/106, 142/159, 39/41; 2010–14: 0/7, -23/-32, 1/-1, 0/-6, 3/2; 2015–19: 6/-18, -5/-35, 0/-5, 4/-20, -1/-30; 2020–24: -37/-42, -22/-24, -35/-40, -12/-23, 5/-12.

A5. Honesty log. Module A: hypothesis commit 10:54 UTC; in-sample commit 11:11 with OOS locked; OOS evaluated once and committed 11:12. Module B: hypothesis and frozen constituent list committed 11:53:57, before any price download; in-sample committed 12:01:40; OOS evaluated once, committed 12:02:08. Module C: hypothesis committed 18:14:13, before any VRP or term-structure series existed; in-sample committed 18:19:02; OOS evaluated once, committed 18:19:38. The sub-period split of B (2005–2012 / 2013–2024) was decided after seeing the by-year table and is labelled descriptive. One bug fixed after A's first in-sample run: factor betas were in mixed units; alphas and t-stats unaffected. 80 variants in total. The practitioner material we were given (an intraday gap-fade strategy matrix and ML-ops code) contained no price or trade data and was used only to motivate Module B's design.

